

Cost is dominated by GPU end-to-end runtime, not tokens

